SHADOW AUDIT

AI Governance Portfolio — Arken Systems

DWP Automated Decision-Making

A governance assessment of the Universal Credit Fraud Detection ML Model and ESA Online Medical Matching Tool, evaluated against the EU AI Act (Annex III), NIST AI RMF, and UK regulatory obligations.

Auditor	Gemma Isaac
Organisation	Arken Systems (fictional)
Subject	Department for Work and Pensions (DWP)
Systems Audited	UC Advance Payment Fraud Detection ML Model; ESA Online Medical Matching Tool
Frameworks	EU AI Act (Annex III), NIST AI RMF, UK GDPR, Equality Act 2010, ICO ADM Guidance
Status	Portfolio artefact — not an official regulatory audit
Date	March 2026
Version	1.0

1. Executive Summary

This audit assesses two AI systems deployed by the Department for Work and Pensions (DWP): the Universal Credit Advance Payment Fraud Detection Machine Learning Model and the Employment and Support Allowance (ESA) Online Medical Matching Tool.

Both systems influence decisions that affect millions of people's access to benefits. These decisions are consequential, often difficult to challenge, and not always visible to those affected.

The audit evaluates both systems against the EU AI Act (Annex III risk classification), the NIST AI Risk Management Framework (AI RMF), and applicable UK obligations under UK GDPR, the Equality Act 2010, and ICO guidance on automated decision-making. The EU AI Act does not apply directly in the UK, but is used here as a benchmark because it provides a more complete model for governing high-risk AI. Where UK law diverges or applies independently, this is noted.

Overall Risk Classification

System	EU AI Act Risk Tier	Audit Finding
UC Advance Payment Fraud Detection ML Model	High-Risk — Annex III, Category 5(b): AI systems used to evaluate eligibility for public assistance benefits	**High**
ESA Online Medical Matching Tool	High-Risk — Annex III, Category 5(b)	**High**

Headline Findings

- Both systems qualify as high-risk under EU AI Act Annex III and trigger the full suite of obligations including conformity assessment, bias monitoring, human oversight, and transparency requirements.
- The fraud detection ML model was deployed covertly between 2020 and 2024 with a documented accuracy rate of only 35%, without public disclosure or adequate fairness analysis across all protected characteristics.
- DWP's own fairness analysis (February 2024, released under FOI) confirmed statistically significant disparities across age, disability, marital status and nationality, yet the department assessed these as not constituting immediate discrimination concerns — a conclusion this audit challenges.
- As of early 2025, DWP had registered only one of an estimated nine to twelve active AI tools on the mandatory Algorithmic Transparency Recording Standard (ATRS), in breach of its own published AI Security Policy.
- No fairness analysis has been conducted for several protected characteristics including race, sex, sexual orientation, religion, pregnancy, and gender reassignment.
- The 'human in the loop' safeguard claimed by DWP is undermined by the volume and pace of casework, raising serious questions about whether human oversight is substantive or nominal.

Auditor's Summary Assessment

Fraud in the benefits system is a real and legitimate concern, and the use of machine learning to identify potentially fraudulent claims is not, in itself, problematic. Used well, these tools can improve consistency and help manage scale. The issue is not whether AI should be used, but whether the governance around these systems is strong enough for the decisions they influence.

Based on the available evidence, there are clear gaps. Both systems appear to have been deployed without full transparency to the people affected, without complete fairness testing across all protected characteristics, and without independent scrutiny. The earlier version of the fraud detection model operated for several years with a reported accuracy rate of 35%, which suggests it generated a high volume of incorrect or unnecessary referrals. The ESA matching tool also produced errors that had direct implications for how claims were handled.

The core concern is that these are not isolated technical issues — they point to weaknesses in how the systems were governed. Taken together, the evidence suggests that both systems would struggle to meet the expectations set out in the EU AI Act for high-risk systems, and that they fall short of the standards DWP has set for itself through its own policies and transparency commitments. For systems operating at this scale, that is a significant concern.

2. System Description

2.1 Universal Credit Advance Payment Fraud Detection ML Model

Universal Credit (UC) advance payments allow claimants to receive an early payment to cover the five-week wait while a new UC claim is processed. The ML model assesses these claims for fraud risk and recommends cases for further human review.

Purpose	Flag UC advance payment claims as potentially fraudulent for human review
Deployment period	July 2020 – present (updated model from 2024)
Scale	Over 780,000 cases processed since deployment
Model type	Machine learning (specific architecture not publicly disclosed)
Training data	Historical claimant data and fraud referrals
Output	Risk score / referral recommendation to caseworker
Human oversight	DWP states a caseworker reviews all flagged cases
Accuracy (v1)	35% correct match rate (2020–2024); 65% required correction by human agent
Transparency	ATRS record: not published as of audit date

2.2 ESA Online Medical Matching Tool

The ESA medical matching tool processes claims for Employment and Support Allowance — the main health-related benefit for people whose capability for work is limited by a condition or disability, claimed by approximately 1.5 million people. When a claimant shares their diagnosis or symptoms, the AI matches the condition to the DWP's incapacity reference guide to register the claim type.

Purpose	Match claimant-reported medical conditions to DWP incapacity reference categories
Deployment period	2020 – present
Scale	Approximately 1.5 million ESA claimants affected
Model type	Large language model (updated 2024)
Output	Condition category match, reviewed by human agent
Accuracy (v1)	35% correct match rate (2020–2024); errors included 'chronic fatigue' matched to 'chronic renal failure'
Accuracy (v2)	87% (based on 360-case sample — sample size queried by researchers)
Transparency	Only DWP AI tool listed on ATRS as of audit date

2.3 Populations Affected

Both systems disproportionately affect individuals who are already in vulnerable circumstances: people with disabilities, those with long-term health conditions, those experiencing unemployment or financial hardship, and recent migrants. These overlapping characteristics are relevant both to the equity analysis and to the adequacy of the human oversight claims.

Universal Credit claimants include a high proportion of younger adults, single parents, people with disabilities, and those in insecure or low-paid work. ESA claimants are, by definition, people whose capacity to work is limited by health conditions — including many with mental health conditions, chronic illness, and disabilities. For these groups, a delayed, withheld, or incorrectly categorised benefit payment can have serious practical consequences.

DWP's own data shows that UC advance payments are disproportionately claimed by people in acute financial difficulty — those for whom the five-week wait is genuinely unmanageable. Applying a model with a 35% accuracy rate to this population, without disclosure or independent oversight, carries risks that go beyond the statistical. The NIST AI RMF's Map function requires that risk identification be grounded in the specific context of deployment, not just the technical characteristics of the model. A governance approach that treats algorithmic errors as uniformly distributed, when the affected population is concentrated among those least able to absorb harm, is not adequate to the actual risk profile of these systems.

3. EU AI Act Risk Classification

The EU AI Act, which entered into force in August 2024, establishes a tiered risk framework for AI systems. Although the Act applies only within the EU and does not bind the UK following Brexit, it represents the most comprehensive regulatory framework for high-risk AI globally and provides the primary benchmark for this audit. Where UK law diverges or applies independently, this is noted.

3.1 Classification

Both DWP systems fall clearly within Annex III of the EU AI Act, which designates the following as high-risk AI:

- Category 5(b): AI systems intended to be used to evaluate the eligibility of natural persons for public assistance benefits and services, including healthcare services, as well as to grant, reduce, revoke, or reclaim such benefits and services.

The UC fraud detection model materially influences whether a claimant is flagged for investigation — a determination that can delay or withhold payment. The ESA medical matching tool determines how a claimant's condition is categorised, which shapes their eligibility pathway. Both systems therefore meet the Annex III threshold.

3.2 Obligations Triggered

Obligation (EU AI Act)	DWP Position (from public sources)	Assessment
Risk management system (Art. 9)	DWP states governance and safeguards exist; limited public documentation	Partial
Data governance and training data (Art. 10)	Trained on historical claims data; no public data card or model card	Non-Compliant
Technical documentation (Art. 11)	Not publicly available; partially disclosed via FOI	Non-Compliant
Logging and record-keeping (Art. 12)	Internal audit trails likely; not independently verifiable	Partial
Transparency to deployers (Art. 13)	Limited; ATRS records largely absent	Non-Compliant
Human oversight (Art. 14)	Caseworkers review flagged cases; adequacy contested	Partial
Accuracy, robustness and cybersecurity (Art. 15)	v1 model: 35% accuracy; v2: 87% on small sample	Partial
Conformity assessment (Art. 43)	No independent third-party assessment evidenced	Non-Compliant
Registration (Art. 51)	ATRS equivalent — one of ~12 tools registered	Non-Compliant

3.3 Jurisdictional Note: UK Position

The EU AI Act does not apply in the UK. The UK government has instead adopted a principles-based, sector-led approach to AI regulation. Relevant UK obligations include:

- UK GDPR Article 22: Rights in relation to automated decision-making and profiling
- Equality Act 2010: Prohibition on indirect discrimination arising from algorithmic bias
- ICO Guidance on Explaining Decisions Made with AI (co-published with the Alan Turing Institute)
- Algorithmic Transparency Recording Standard (ATRS): Government-mandated disclosure framework
- Public Sector Equality Duty (s.149 Equality Act 2010): Positive obligation to consider equality impacts

The EU AI Act is used in this audit as a best-practice governance benchmark. Where a gap exists under EU AI Act standards, the audit also considers whether an equivalent gap exists under UK law.

4. NIST AI RMF Assessment

The NIST AI Risk Management Framework (AI RMF 1.0, 2023) provides a voluntary but internationally recognised structure for managing AI risk across four core functions: Govern, Map, Measure, and Manage. This section assesses DWP's observable practices against each function.

4.1 GOVERN

Policies, processes, accountability, and culture for AI risk management

Area	DWP Observable Practice	Finding
AI governance policy	DWP AI Security Policy published (2025); Digital Design Authority exists as governance board	Partial
Senior accountability	Director General for Fraud identified; ministerial oversight in place	Compliant
Third-party oversight	No evidence of independent audit; internal governance only	Non-Compliant
Staff training on AI risk	Not evidenced in public documentation	Non-Compliant

4.2 MAP

Identifying and categorising AI risks in context

Area	DWP Observable Practice	Finding
Risk identification	NAO flagged inherent bias risk; DWP acknowledged but disputed significance	Partial
Affected populations	Recognised in Equality Analysis; scope of characteristics limited	Partial
Intended vs actual use	v1 model deployed covertly 2020–2024; limited public purpose statement	Non-Compliant
Contextual risk factors	Vulnerability of claimant population not fully reflected in risk weighting	Non-Compliant

4.3 MEASURE

Analysing and assessing AI risks quantitatively and qualitatively

Area	DWP Observable Practice	Finding

Accuracy measurement	v1: 35% (acknowledged); v2: 87% on 360-case sample — methodology queried	**Partial**
Bias and fairness analysis	February 2024 analysis conducted; significant protected characteristics excluded	**Partial**
Independent measurement	No third-party fairness audit evidenced	**Non-Compliant**
Ongoing monitoring	Enhanced Review Team and TCR agents review flagged cases; systemic monitoring not evidenced	**Partial**

4.4 MANAGE

Prioritising and treating identified AI risks

Area	DWP Observable Practice	Finding
Human override mechanisms	Caseworker review stated; adequacy at scale contested	**Partial**
Incident response	Not documented in public sources	**Non-Compliant**
Redress for affected individuals	Formal appeals process exists for benefit decisions; AI component not separately challengeable	**Partial**
Transparency to affected individuals	No notification to claimants that AI was used in their case	**Non-Compliant**

Across the four functions, a consistent pattern emerges. DWP's governance infrastructure is most developed at the policy and accountability layer — senior responsibility is assigned, a governance board exists, and internal policies are published. But as the assessment moves closer to operational reality, the picture becomes less clear. The Map function shows gaps in how risk has been identified in context, particularly in relation to the specific vulnerabilities of affected populations. The Measure function produced a fairness analysis that covers only four of the nine Equality Act protected characteristics, and whose conclusions have been challenged by independent researchers.

The weakest area is the Manage function. Human override mechanisms are structurally constrained by caseload volume; incident response processes are not documented in public sources; and affected individuals have no practical way to know that AI was involved in their case, let alone to challenge it. These are the points where governance either protects people or fails them. Based on the available evidence, there are real gaps at this level.

The NIST AI RMF is a voluntary framework, and DWP has not formally adopted it. Its use here is analytical. But its four functions represent a broadly accepted standard for responsible AI deployment, and the pattern of gaps identified is consistent across all the frameworks applied in this audit — which suggests the issues are structural rather than incidental.

5. Governance Gap Analysis

This section examines the five most significant governance failures identified in this audit. Each gap is assessed for its severity, its regulatory significance, and its potential impact on affected individuals.

Gap 1 — Transparency Failure: ATRS Non-Compliance

Severity	**High**
Regulatory basis	Algorithmic Transparency Recording Standard (mandatory); DWP AI Security Policy s.11; GDPR Art. 13/14 transparency obligations
Finding	Despite operating between nine and twelve AI tools, DWP had registered only one on the ATRS as of early 2025. The ESA medical matching tool is listed; the UC fraud detection model is not. This is non-compliant with the ATRS mandatory regime and DWP's own AI Security Policy.

Transparency is not just a procedural requirement. It is what allows any other form of accountability to exist in practice. If it is not clear which systems are in use, what they do, and how they are governed, neither affected individuals nor oversight bodies are in a position to question or challenge them. The ATRS was designed to address this by providing a public record of algorithmic tools used in public sector decision-making, with mandatory disclosure requirements that DWP's own AI Security Policy confirms apply to its systems.

As of early 2025, DWP had registered one tool — the ESA Online Medical Matching Tool — while operating between nine and twelve AI systems. The UC fraud detection model, which processed over 780,000 cases, was not registered. A 2025 report by researchers at the University of Edinburgh, analysing Freedom of Information requests submitted to DWP, found that the department had been opaque and inconsistent in its disclosures, with some requesters receiving documents only after ICO complaints. The researchers called for DWP to make all relevant documents available through the ATRS rather than requiring civil society organisations to pursue them case by case.

The practical consequence is that the public, parliamentarians, researchers, and the people directly affected by these systems cannot independently assess whether they are operating fairly or within the law. This is not simply a gap in documentation. It limits the ability of others to identify problems or hold the system to account. A department that describes itself as committed to the safe, ethical, and value-driven adoption of AI, while withholding basic transparency records, is not meeting either the spirit or the letter of its own stated standards.

Gap 2 — Bias Methodology: Incomplete Protected Characteristics

Severity	**High**
Regulatory basis	Equality Act 2010 (all nine protected characteristics); Public Sector Equality Duty; EU AI Act Art. 10 (data governance)
Finding	The February 2024 fairness analysis covered age, disability, marital status, and nationality. No analysis has been conducted for race, sex, sexual orientation, religion, pregnancy, maternity, or gender reassignment — five of the nine Equality Act 2010 protected characteristics.

Partial fairness testing is not the same as demonstrating fairness. The February 2024 analysis covered some protected characteristics but not all those required under the Equality Act 2010. This means there is no evidence available for how the system performs across

several legally relevant groups — including race, sex, sexual orientation, religion, pregnancy, maternity, and gender reassignment. The absence of measurement is not evidence of absence.

Where disparities were found, they were assessed as not constituting immediate discrimination concerns. That conclusion depends heavily on how concern is defined — a standard that is not clearly set out, independently defined, or legally grounded. DWP's framing of the results was produced by the same organisation responsible for deploying the system, which limits the weight that can reasonably be placed on it.

There is also a question about what the redactions conceal. Key findings were withheld on the basis that disclosure might allow fraudsters to game the system. That justification has not been tested by the ICO, and civil society organisations have questioned it. Where the characteristics most likely to reveal uneven treatment — such as nationality — were examined, the results were partially obscured. A fairness analysis that cannot be independently scrutinised cannot provide meaningful public assurance. Without full coverage and clear methodology, it is difficult to treat the analysis as providing the assurance that is being claimed for it.

Gap 3 — Covert Deployment and Accuracy Failure

Severity	**High**
Regulatory basis	UK GDPR Art. 22 (right to explanation); ICO ADM guidance; EU AI Act Art. 15 (accuracy and robustness standards)
Finding	The UC fraud detection model operated from 2020 to 2024 with a 35% accuracy rate, meaning 65% of referrals required human correction. The ESA medical matching tool similarly produced serious condition mismatches. Both were deployed without public disclosure or notification to affected claimants.

The 35% accuracy figure is significant, but the governance concern is not just the number itself. It means that a large proportion of cases flagged by the system required correction by a caseworker. In practice, this likely increased administrative load rather than reducing it, while also introducing friction for claimants in genuine financial need. The ESA medical matching tool produced errors of a different kind but comparable gravity — including condition mismatches with direct implications for how claims were categorised and handled.

The governance issue is not primarily that these systems had limitations. Early-stage models often do, and refinement over time is expected. The problem is that both appear to have been used without clear disclosure to the people affected. Claimants who were flagged by the fraud detection model had no way of knowing an algorithm had contributed to their case being reviewed. Those whose ESA applications were mismatched had no way to identify or challenge the AI component specifically. This limits the ability of individuals to understand or seek review of decisions that involve automated processing — which is directly relevant to rights under UK GDPR Article 22.

From a governance perspective, the combination of limited transparency and low initial accuracy is where the concern sits. Systems operating at this scale, affecting people in vulnerable circumstances, need to be visible and subject to challenge. The secrecy was not incidental — it removed the conditions under which external scrutiny and correction could take place.

Gap 4 — Self-Assessment and the Adequacy of Internal Review

Severity	**Medium**
Regulatory basis	EU AI Act Art. 43 (conformity assessment — third-party requirement for high-risk systems); ICO AI strategy 2025 (independent assurance)
Finding	DWP's fairness analysis was conducted internally and concluded the identified disparities did not translate to immediate discrimination concerns. No independent third-party audit of either system has been evidenced. For systems operating at this scale and consequence, self-assessment is structurally insufficient.

The fairness analysis appears to have been conducted internally and reached relatively reassuring conclusions. That is not unusual, but it does raise a question about independence. Where the same organisation is responsible for deploying a system and assessing its impact, there is a structural risk that issues are interpreted conservatively. DWP's Director General for Fraud acknowledged in a select committee session that the algorithms contain biases, and offered the view that some bias is necessary to identify fraudulent claims. That position is not technically unreasonable, but it illustrates the problem: the question of whether a given bias is acceptable under law is not one that can be settled by the department whose performance targets depend on the model's outputs.

DWP has invested £70 million in advanced analytics between 2022 and 2025 and projects significant savings from these systems. For systems operating at this scale and consequence, external scrutiny becomes more important, not less. Independent review does not remove all risk, but it provides a different perspective and helps test conclusions that internal processes may be inclined to validate. The ICO's 2025 AI strategy explicitly identifies high-stakes ADM in central government as a priority area. This audit's findings suggest DWP's systems would merit that attention.

Gap 5 — Human Oversight: Substantive or Nominal?

Severity	**Medium**
Regulatory basis	EU AI Act Art. 14 (human oversight); UK GDPR Art. 22 (meaningful human review requirement); ICO ADM guidance
Finding	DWP consistently states that 'a caseworker will always look at all available information to make a decision.' However, each work coach is responsible for over 100 claimants, with appointments as short as ten minutes. At this volume, the cognitive capacity for meaningful override of AI recommendations is structurally constrained.

DWP states that a caseworker reviews decisions, which is intended to provide a safeguard. Whether that safeguard is effective depends on how it works in practice.

Each work coach is responsible for over 100 claimants, with appointments lasting as little as ten minutes. In that context, it may not always be possible to fully examine an AI-generated recommendation — particularly where the system output is presented as a risk signal rather than an input to be weighed. Research on automation bias shows that human reviewers under time pressure tend to accept algorithmic recommendations at higher rates than they would reach the same conclusion independently. This is not a failure of individual caseworkers. It is a predictable consequence of deploying AI recommendations in conditions that constrain deliberate, independent review.

This does not mean that oversight is absent, but it raises a question about how substantive it is. Meaningful human review — as required under UK GDPR Article 22 and the EU AI Act Article 14 — means having both the authority and the practical capacity to question or override the system where needed. Whether current caseworker conditions meet that standard is a question that has not been formally assessed and that this audit would recommend examining.

6. Regulatory Framework Comparison

This section maps the governance gaps identified in Section 5 against the applicable regulatory frameworks, noting where obligations arise under UK law independently of the EU AI Act.

Gap	EU AI Act	UK GDPR	Equality Act	ATRS/ICO
ATRS Non-Compliance	Art. 51 (registration)	Art. 13/14 (transparency)	s.149 PSED	ATRS mandatory — direct breach
Incomplete bias testing	Art. 10 (data governance)	Art. 5(1)(a) fairness	All 9 protected characteristics	ICO ADM guidance
Covert deployment / accuracy	Art. 15 (accuracy), Art. 13	Art. 22 (right to explanation)	Indirect discrimination risk	ICO AI strategy
Self-assessment only	Art. 43 (conformity assessment)	N/A (no direct equivalent)	PSED — adequate evidence required	ICO independent assurance guidance
Nominal human oversight	Art. 14 (human oversight)	Art. 22 (meaningful review)	N/A	ICO ADM — meaningful not token

6.1 UK vs EU Divergence

The UK has taken a more principles-based approach to AI governance rather than adopting a binding high-risk classification system similar to the EU AI Act. This does not mean that no obligations apply. UK GDPR, the Equality Act, the ATRS, and sector-specific requirements all remain relevant. But the structure is different.

In practice, accountability is spread across multiple mechanisms rather than being concentrated in a single framework. That can make it harder to identify where responsibility sits and how gaps should be addressed. The ATRS is mandatory but lacks a clearly defined enforcement mechanism. UK GDPR Article 22 rights depend on individuals knowing AI was used — which covert deployment structurally prevents. The ICO has identified high-stakes ADM in central government as an enforcement priority but has not yet acted in relation to DWP's systems.

The result is not necessarily weaker protection in every case, but it does mean that oversight is less standardised and, in some cases, less predictable. For the people affected by these systems, the practical difference matters. The absence of an equivalent to the EU AI Act's conformity assessment requirement means there is no mandatory external check before high-risk systems go live. That gap is relevant to how the issues identified in this audit arose, and to what would be needed to prevent similar issues in future.

7. Remediation Recommendations

The following recommendations are addressed to DWP as the deploying organisation. They are ordered by priority, with immediate actions first.

#	Recommendation	Rationale	Priority
1	**Publish full ATRS records for all active AI tools**	Immediate compliance with mandatory ATRS obligations and DWP's own AI Security Policy. Transparency is a prerequisite for accountability.	Immediate
2	**Commission independent third-party fairness audit across all nine Equality Act protected characteristics**	Self-assessment has demonstrably failed to identify the full scope of potential bias. Independent audit is required for systems of this consequence.	Immediate
3	**Issue retrospective notification to claimants affected by v1 UC fraud model (2020–2024)**	UK GDPR transparency obligations were not met during covert deployment. Claimants have a right to know AI was used in decisions affecting them.	Short-term
4	**Conduct formal human oversight capacity review**	Caseworker caseloads and appointment durations must be assessed against what constitutes meaningful review of AI recommendations under GDPR Art. 22.	Short-term
5	**Publish model cards for both systems**	Model cards documenting training data, intended use, limitations, and known biases are standard practice for responsible AI deployment and would address multiple transparency gaps.	Medium-term
6	**Establish an independent AI ethics advisory function**	An external advisory body with claimant representation would provide ongoing scrutiny independent of internal governance structures.	Medium-term
7	**Adopt the EU AI Act Annex III framework voluntarily as a governance standard**	In the absence of equivalent UK legislation, voluntarily adopting the EU AI Act's high-risk obligations would provide a clear, internationally recognised benchmark and reduce legal risk.	Long-term

8. Audit Limitations

Transparency about the constraints of this audit is itself a governance principle. The following limitations are material to interpreting the findings.

8.1 Source Constraints

- This audit relies entirely on publicly available sources: parliamentary evidence, FOI disclosures, NAO reports, ICO publications, academic research, and journalism. Key documents — including the full fairness analysis, DPIAs, and model technical documentation — remain partially or wholly redacted.
- The opacity of DWP's systems is itself a finding. The difficulty of conducting an independent assessment from public sources reflects the transparency gaps identified in Section 5.
- This is not an official regulatory audit. It does not carry the legal authority of an ICO investigation or parliamentary scrutiny exercise.

8.2 Methodology

- Framework assessment (EU AI Act obligations, NIST AI RMF functions) is based on publicly documented DWP practices. Where documentation is absent, this audit treats absence as a negative indicator — the burden of demonstrating compliance lies with the deploying organisation.
- The 'Partial' classification is used where some evidence of compliance exists but the evidence is insufficient to establish full compliance.

8.3 What This Audit Cannot Determine

- Whether the current (post-2024) UC fraud detection model has improved fairness outcomes across all protected characteristics.
- The specific decision weights or features used in either model.
- Whether internal DPIAs and governance processes adequately address the risks identified — these documents have not been made public.
- The precise causal relationship between algorithmic flags and final benefit decisions.

8.4 Auditor's Position

This audit was conducted by Gemma Isaac as a portfolio artefact. It does not constitute legal advice, a regulatory determination, or an official audit under any statutory framework. The conclusions are analytical judgements based on evidence available in the public domain and should not be treated as legal findings.

The author has no institutional affiliation with DWP, the ICO, or any organisation referenced in this document. Sources have been cited in good faith. Where documents were partially redacted, this audit has noted the redaction rather than speculating about what was withheld.

Some of the sources cited — including reports from Amnesty International, the Public Law Project, and Novara Media — have campaigning positions on DWP's use of AI. This audit has also drawn on DWP's own publications, NAO reports, and parliamentary evidence. Where sources agree on a specific factual point, that point has been treated as established. Where interpretation differs, this audit has sought to reason from the evidence rather than adopt any single position. The reader is encouraged to consult the primary sources in Section 9 and form their own view.

9. Sources and References

All sources used in this audit are publicly available. Where documents were obtained via Freedom of Information requests, the requesting organisation is noted.

Primary Sources

- Department for Work and Pensions. Artificial Intelligence Security Policy. GOV.UK, May 2025.
- Department for Work and Pensions. Annual Report and Accounts 2023–2024.
- DWP Advances Model Fairness Analysis (February 2024). Released under FOI to the Public Law Project.
- National Audit Office. Fraud and Error in the Benefits System (2023). HC [reference to be confirmed].
- UK Parliament. Select Committee Evidence Session on AI in Welfare. October 2023.
- Algorithmic Transparency Recording Standard (ATRS) — DWP entry (ESA Online Medical Matching Tool). GOV.UK.

Secondary Sources and Reports

- Amnesty International. Social Insecurity: The Devastating Human Rights Impact of Social Security System Failures in the UK. 2025.
- Currie, M. & Spring, A. Algorithmic Accountability in the UK. University of Edinburgh / ERA, July 2025.
- Public Law Project. Tracking Automated Government (TAG) Register. publiclawproject.org.uk.
- Privacy International. Written Evidence to Parliament on DWP Automated Decision-Making.
- Tony Blair Institute. Governing in the Age of AI: Reimagining the UK Department for Work and Pensions. July 2024.

Regulatory Frameworks

- EU AI Act (Regulation (EU) 2024/1689). Official Journal of the European Union.
- NIST AI Risk Management Framework (AI RMF 1.0). National Institute of Standards and Technology, 2023.
- UK GDPR. Retained EU Law (incorporated via European Union (Withdrawal) Act 2018).
- Equality Act 2010 (UK). c.15.
- ICO. Guidance on Automated Decision-Making and Profiling (under review following Data (Use and Access) Act 2025).
- ICO & Alan Turing Institute. Explaining Decisions Made with AI.
- ICO. AI and Biometrics Strategy 2025–26.

Journalism and Commentary

- Computer Weekly. DWP Fairness Analysis Reveals Bias in AI Fraud Detection System. December 2024.
- Big Issue. Warning After DWP Reveals How It Uses AI to Make Decisions on Benefits. February 2025.

- Novara Media. How Labour Is Using Biased AI to Determine Benefit Claims. April 2025.
- Tech Monitor. UK's AI System for Welfare Fraud Detection Faces Criticism Over Bias and Transparency. December 2024.
- Public Technology. DWP Working to Finalise Publication of More Algorithm Details. April 2025.
- Government Transformation. DWP Turns to Tech to Tackle Benefit Fraud — But Integration and Fairness Issues Remain. October 2025.

This document is a portfolio artefact produced for the purpose of demonstrating AI governance competence. It does not constitute legal advice or an official regulatory determination. All source attributions are made in good faith based on publicly available information.